

Distributed Databases

And their relevance for NOSQL

NOSQL Goals

Recall the NOSQL Goals:

- Schemaless: Allow flexible (even runtime) schema definition [data structure]
- Reliability / availability: Keep delivering service even if its software or hardware components fail [recovery]
- Scalability: Continuously evolve to support a growing amount of tasks [distribution]
- Efficiency: How well the system performs, usually measured in terms of response *time* (latency) and *throughput* (bandwidth) [distribution]

NOSQL Goals

Recall the NOSQL Goals:

- Schemaless: Allow flexible (even runtime) schema definition [data structure]
- Reliability / availability: Keep delivering service even if its software or hardware components fail [distribution]
- Scalability: Continuously evolve to support a growing amount of tasks [distribution]
- Efficiency: How well the system performs, usually measured in terms of response *time* (latency) and *throughput* (bandwidth) [distribution]

NOSQL Goals

- Reliability / availability: Keep delivering service even if its software or hardware components fail [distribution]
- Scalability: Continuously evolve to support a growing amount of tasks [distribution]
- Efficiency: How well the system performs, usually measured in terms of response *time* (latency) and *throughput* (bandwidth) [distribution]

NOSQL Goals

- Reliability / availability: Keep delivering service even if its software or hardware components fail [distribution]
- Scalability: Continuously evolve to support a growing amount of tasks [distribution]
- Efficiency: How well the system performs, usually measured in terms of response *time* (latency) and *throughput* (bandwidth) [distribution]

Most NOSQL goals can be achieved by means of distribution!

Activity: Why Distribution?

- Objective: Recognize the benefits of distributing data

- Tasks:

1. (15') By pairs, answer the following questions:

- a) How long would it take to read 1TB with sequential access (fig. a)? (in secs)

- Can you identify any additional drawback to be considered?

- b) How long would it take to read 1TB with parallel access (fig. b)? Assume 100 disks on the same machine with shared-memory and infinite CPU capacity.

- Can you identify any additional drawback to be considered?

- c) How long would it take to read 1TB with distributed access (fig. c)? Assume 100 shared-nothing machines in a star-shape LAN in a single rack where all data is sent to the centre.

- Can you identify any additional drawback to be considered?


- d) Now, repeat the exercise considering a single random access. What changes?

2. (5') Discussion


Type	Latency	Bandwidth
Disk	$\approx 5 \times 10^{-3}$ s (5 millise.);	At best 100 MB/s
LAN	$\approx 1 - 2 \times 10^{-3}$ s (1-2 millise.);	≈ 1 GB/s (single rack); ≈ 100 MB/s (switched);
Internet	Highly variable. Typ. 10-100 ms.;	Highly variable. Typ. a few MB/s.;

Bottom line (1): it is approx. one order of magnitude faster to exchange main memory data between 2 machines in a data center, that to read on the disk.


Bottom line (2): exchanging through the Internet is slow and unreliable with respect to LANs.



a. Single CPU, single disk



b. Parallel read: single CPU, many disks



c. Distributed reads: an extendible set of servers

Activity: Why Distribution?

- Objective: Recognize the benefits of distributing data
- Tasks:
 1. (15') By pairs, answer the following questions:
 - a) How long would it take to read 1TB with sequential access (fig. a)? (in secs)
 - Can you identify any additional drawback to be considered?
 - b) How long would it take to read 1TB with parallel access (fig. b)? Assume 100 disks on the same machine with shared-memory and infinite CPU capacity.
 - Can you identify any additional drawback to be considered?
 - c) How long would it take to read 1TB with distributed access (fig. c)? Assume 100 shared-nothing machines in a star-shape LAN in a single rack where all data is sent to the centre.
 - Can you identify any additional drawback to be considered?
 - d) Now, repeat the exercise considering a single random access. What changes?
 2. (5') Discussion



Conclusions:

- Disk transfer rate is a bottleneck for batch processing of large scale datasets; parallelization and distribution can help to eliminate this bottleneck
- Disk seek time is a bottleneck for transactional applications that submit a high rate of random accesses; replication, distribution of writes and distribution of reads make such applications scalable
- When possible, data should be accessed where they are stored (or near) to avoid costly data exchange over the network

Definition and Distributed Architectures

Distributed Database

- A distributed database (DDB) is a database where data management is distributed over several nodes in a network.
 - Each node is a database itself
 - Potential heterogeneity
 - Nodes communicate through the network



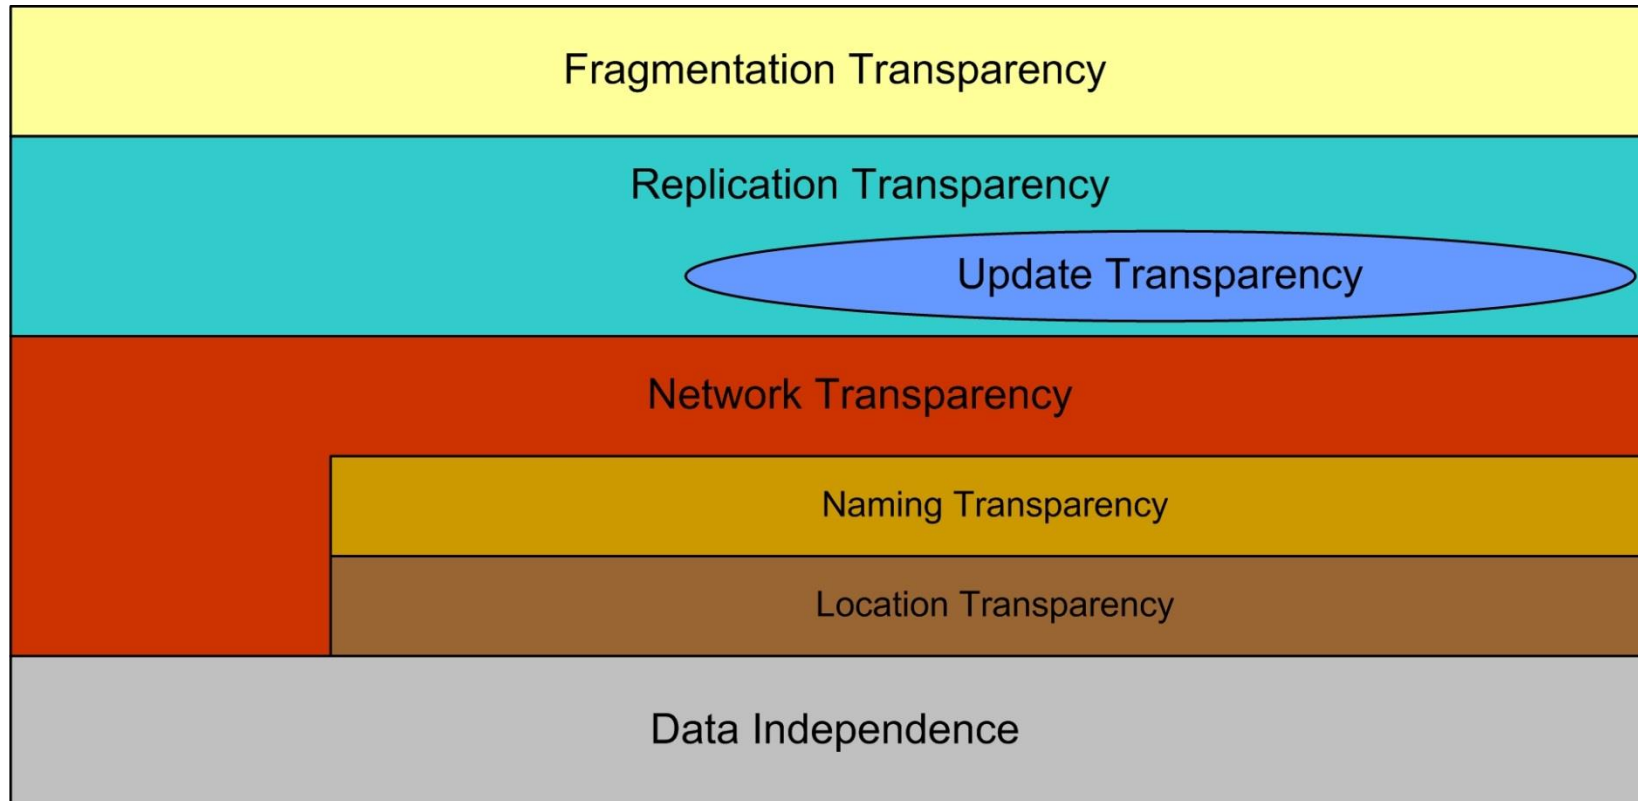
Distributed Architectures

- Main objective: hide implementation (i.e., physical) details to the users who should not know they are dealing with a distributed system
 - Network transparency
 - Data access must be independent regardless where data is stored
 - Each data object must have a unique name
 - Replication transparency
 - The user must not be aware of the existing replicas
 - Fragmentation transparency
 - The user must not be aware of partitioning

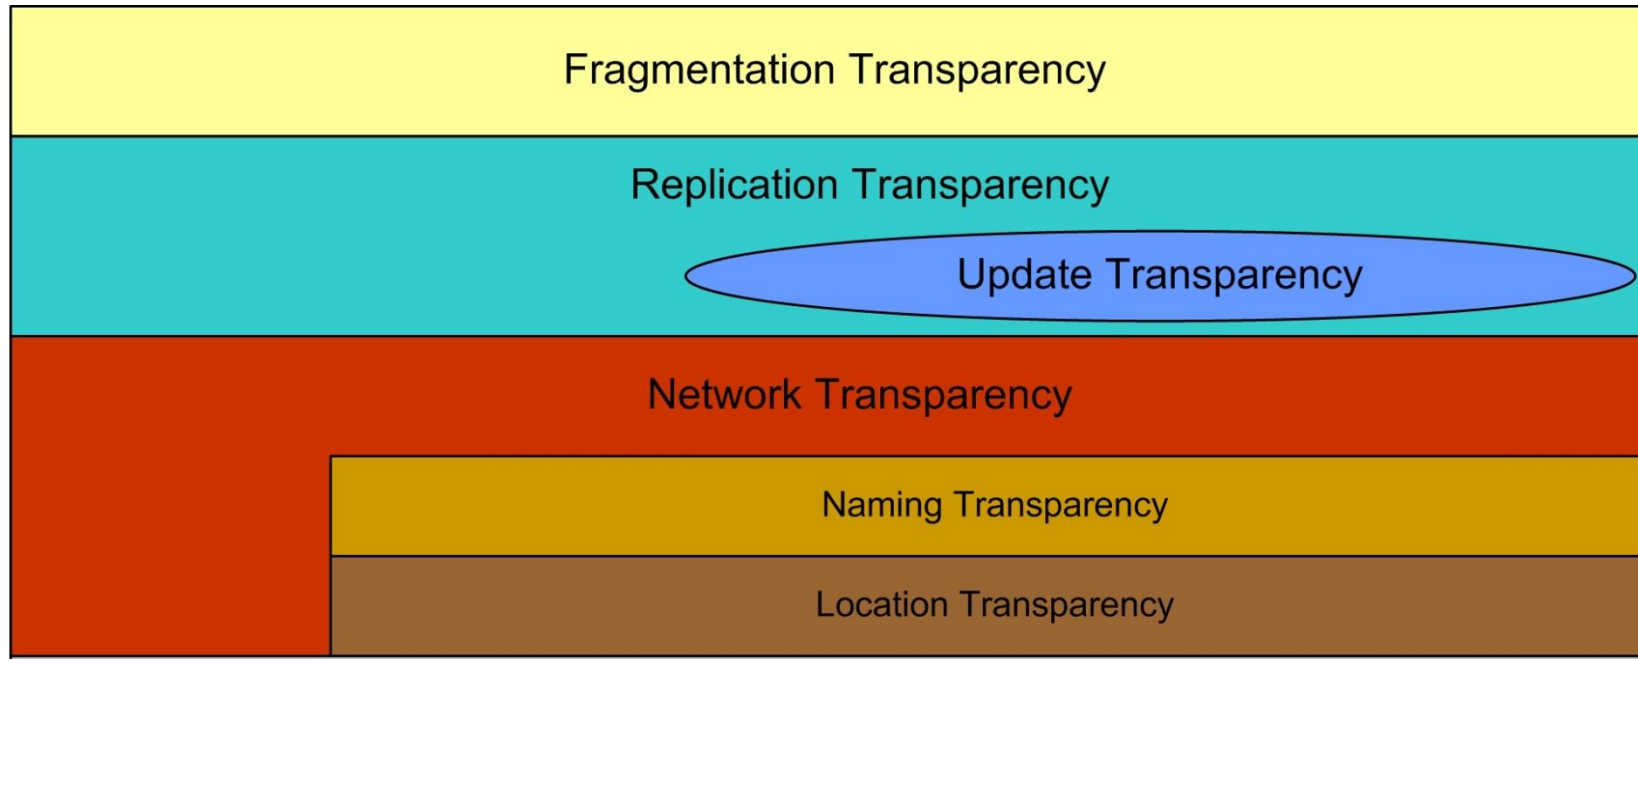
[**Desiderata:** Data independency at the logical and physical level must be guaranteed to avoid accessing directly to physical structures]

- As seen in centralized DBs (ANSI SPARC)

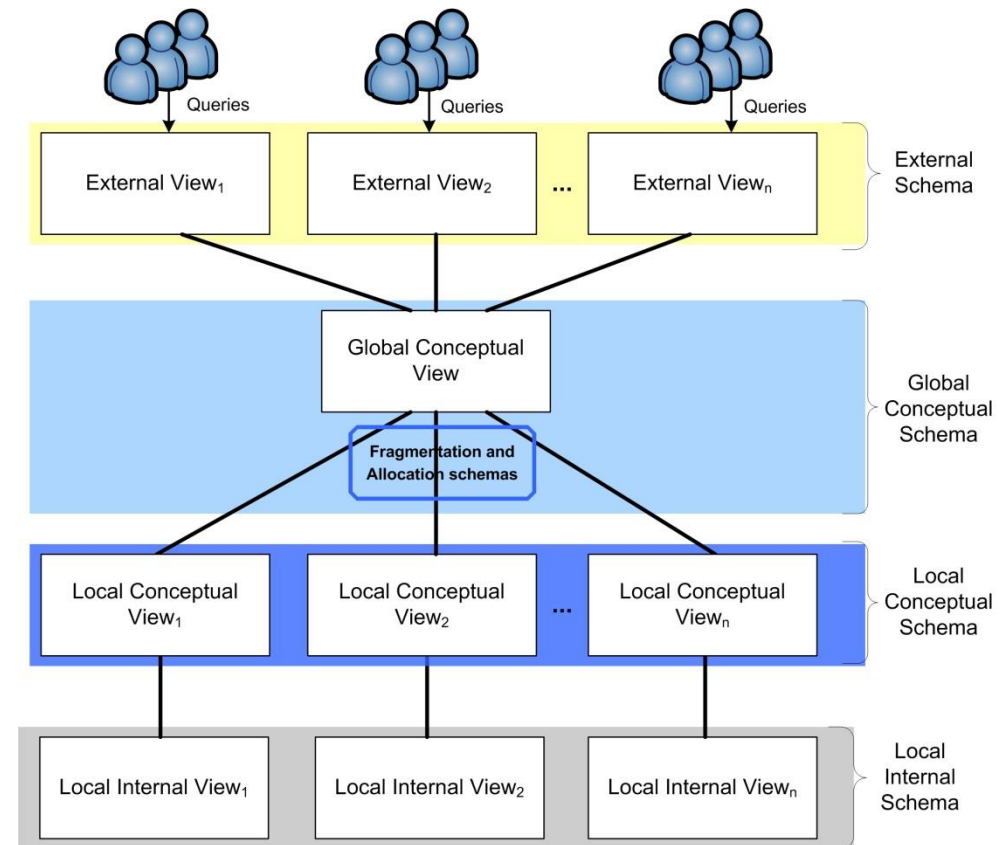

Distributed Architectures



Distributed Architectures



Extended ANSI-SPARC Architecture



- Global catalog
 - Mappings between ESs – GCS and GCS – LCSs
- Each node has a local catalog
 - Mappings between LCS_i – IS_i

Distributed DBMS Architecture

A distributed database has one global layer but many local layers (or nodes). Each node stores and process data locally

Query Manager: unfolds views, checks permissions and constraints and generates (and optimizes) a physical access plan (procedural) to data in disk from the declarative SQL query

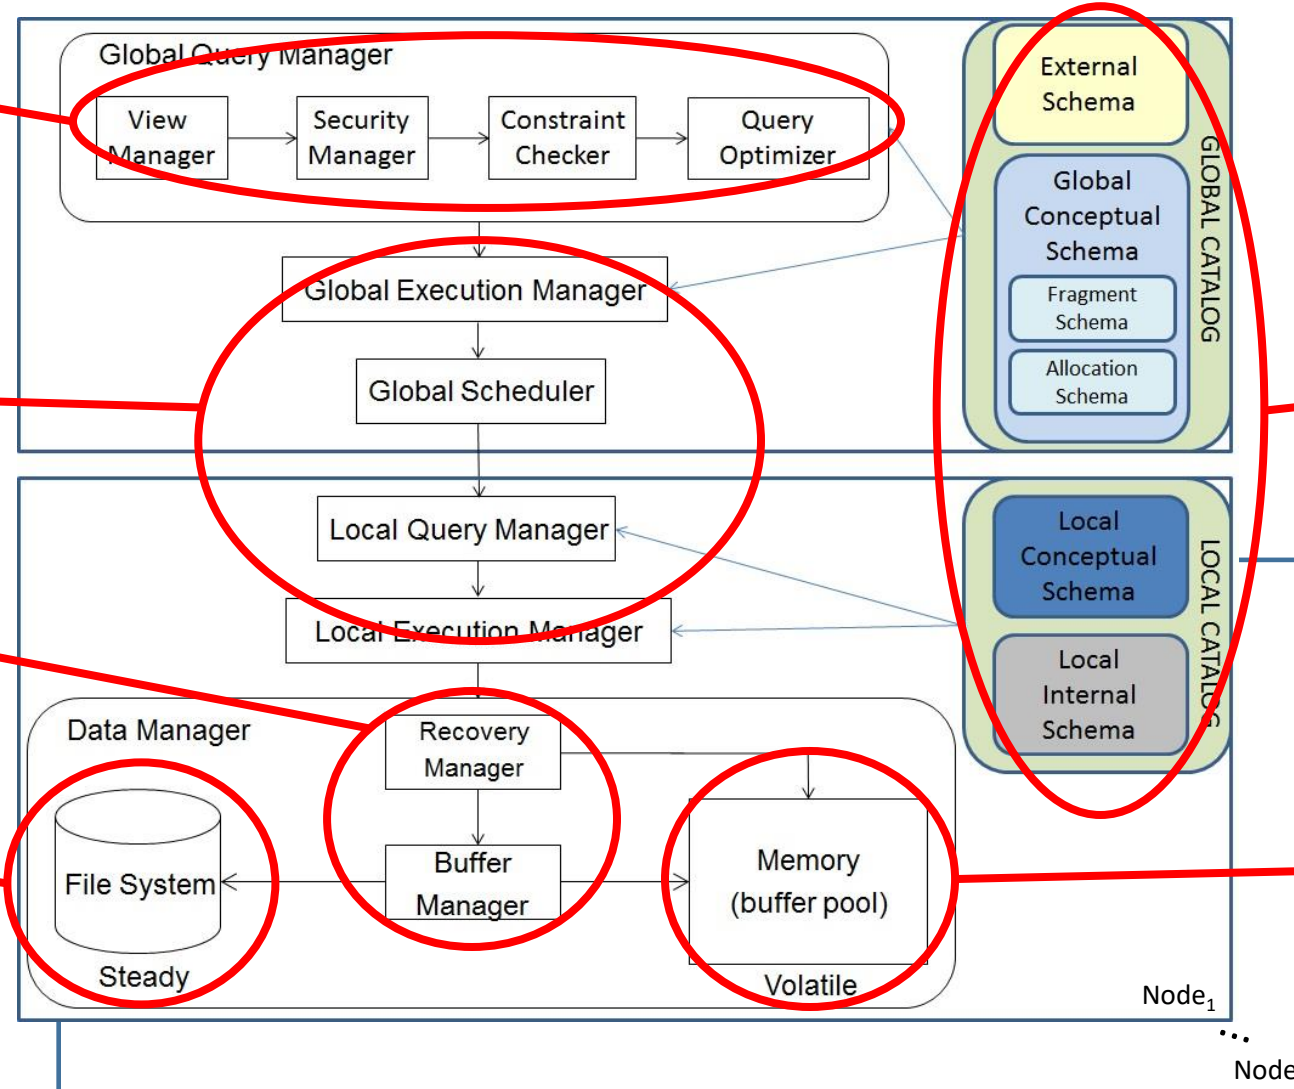
Concurrency manager and scheduler: guarantees CI (from ACID) and schedules the execution of subqueries over the distributed fragments. They are responsible for bridging between the global layer and each node

Recovery manager: interacts with the buffer and memory to guarantee AD (from ACID)

Disk: physical data storage (persistence)

Global and local catalog: needed to make distribution transparent to the end-user

Main memory: memory assigned to the database in that node



Challenges

Principles of Distributed Databases




Challenges in Distributed Databases

- I. Distributed DB design
 - Node distribution
 - Data fragments
 - Data allocation (replication)
- II. Distributed DB catalog
 - Fragmentation trade-off: Where to place the DB catalog
 - Global or local for each node
 - Centralized in a single node or distributed
 - Single-copy vs. Multi-copy
- III. Distributed query processing
 - Data distribution / replication
 - Communication overhead
- IV. Distributed transaction management
 - How to enforce the ACID properties
 - Replication trade-off: Queries vs. Data consistency between replicas (updates)
 - Distributed recovery system
 - Distributed concurrency control system
- V. Security issues
 - Network security

Challenges in Data Distribution

A	B	C	D	E

Conceptual View




Physical View

Challenges in Data Distribution

Conceptual View

A	B	C	D	E

? How to physically store this table in a distributed manner (making it transparent for the user)?



Physical View

Challenge I

Database Design: Fragmentation and Allocation (Replication)

Challenge I: DDB Design

- Given a DB and its workload, how should the DB be split and allocated to sites as to optimize certain objective functions
 - For example, minimize resource consumption for query processing
- Two main issues:
 - Data fragmentation
 - Data allocation (data replication)

Activity: Fragmentation Strategies

- Objective: Understand the principles behind fragmentation
- Tasks:
 1. (10') By pairs, answer the following questions:
 - I. Briefly explain which fragmentation strategy has been applied for the database below.
 - II. Is this fragmentation strategy complete and disjoint? Can we reconstruct the global relations? Explicit the algebraic operation you would use.
 2. (5') Discussion

Global Relations

Kids(kidId, name, address, age)

Toys(toyId, name, price)

Requests(kidId, toyId, willingness)

Note that requests(kidId) is a foreign key to kids(kidId) and similarly, requests(toyId) refers to toys(toyId).

Fragments

K1= Kids[kidId, name]

K2= Kids[kidId, address, age]

T1= Toys(price >= 150)

T2= Toys(price < 150)


R1 = Requests ⋈ T1

R2 = Requests ⋈ T2

Challenge I: Data Fragments

Conceptual View

A	B	C	D	E




Physical View

Challenge I: Data Fragments

Conceptual View


A	B	C	D	E
Apply a fragmentation strategy				



Physical View

Challenge I: Data Fragments

Conceptual View



Physical View


Secondary servers

Challenge I: Data Allocation

- Given a set of fragments, a set of sites on which a number of applications are running, **allocate** each fragment such that some optimization criterion is met (subject to certain constraints)
- It is known to be a NP-hard problem
 - The optimal solution depends on many factors
 - Location in which the query originates
 - The query processing strategies (e.g., join methods)
 - Furthermore, in a dynamic environment the workload and access pattern may change
- Most advanced approaches build *cost models* and any optimization algorithm can be adapted to solve it
 - Simply the problema with assumptions (e.g., only communication cost considered)
 - Heuristics are also available: (e.g., best-fit for non-replicated fragments)
 - Sub-optimal solutions (i.e., applied per fragment individually)

Challenge I: Data Allocation



Conceptual View



Physical View

Challenge I: Data Allocation

Conceptual View



Physical View


Challenge I: Data Allocation

Conceptual View

F2

F3

Apply an
allocation
strategy




Physical View

Challenge I: Data Allocation

Conceptual View

F3 Apply an allocation strategy




Physical View

Challenge I: Data Allocation

Conceptual View


Apply an
allocation
strategy



Physical View

Challenge I: Data Allocation

Conceptual View




Physical View

Challenge I: Data Allocation

Conceptual View


If a fragment is placed in more than one server, then, we are **replicating** it




Physical View

Challenge I: Data Replication Management


- Replicating fragments improves the system throughput but raises some other issues:
 - Consistency
 - Update performance
- Most used replication protocols
 - Eager – Lazy replication
 - Primary – Secondary versioning




a) Eager replication with primary copy



b) Lazy replication with primary copy
(a.k.a Master-Slave replication)



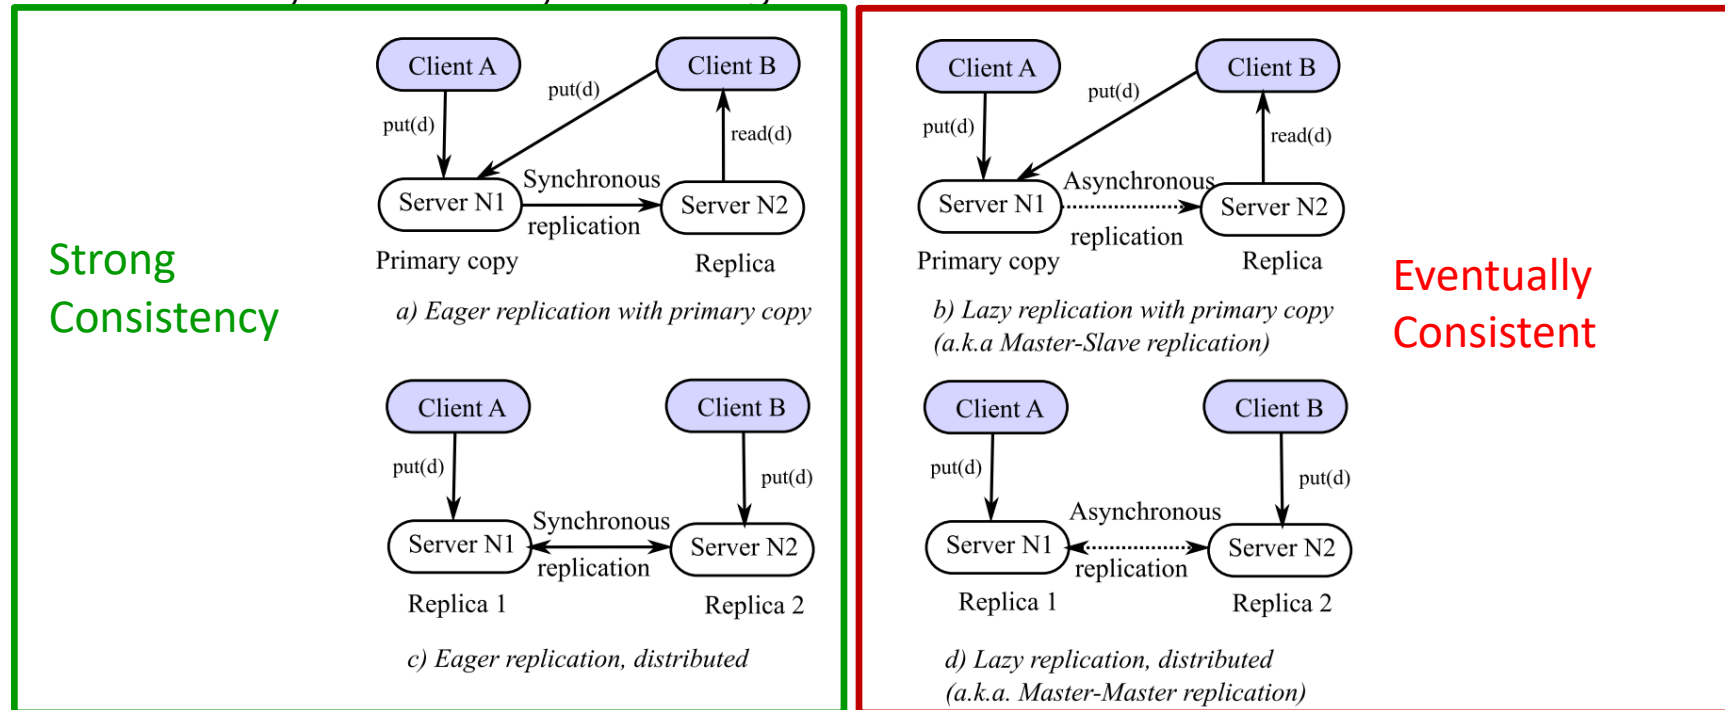
c) Eager replication, distributed



d) Lazy replication, distributed
(a.k.a. Master-Master replication)

Challenge I: Data Replication Management

- Replicating fragments improves the system throughput but raises some other issues:
 - Consistency
 - Update performance
- Most used replication protocols
 - Eager – Lazy replication
 - Primary – Secondary versioning



Activity: Data Replication Issues

- *Objective: Understand the consequences behind each data replication strategy*
- *Tasks:*
 1. (10') *By pairs, answer the following questions:*
 - I. *Discuss the questions below with your peer*
 - II. *What is the most important feature for each scenario?*
 2. (5') *Discussion*
- You are a customer using an e-commerce based on heavy replication (e.g., Amazon):
 - a) Show a database replication strategy (e.g., sketch it) where:
 1. You buy an item, but this item does not appear in your basket.
 2. You reload the page: the item appears.What happened?
 - b) Show a database replication strategy (e.g., sketch it) where:
 1. You delete an item from your command, and add another one: the basket shows both items.What happened?
Will the situation change if you reload the page?

Challenge II


The Global Catalog



Challenge II: Global Catalog


A	B	C	D	E

Conceptual View




Physical View

Challenge II: Global Catalog




Challenge II: Global Catalog



Conceptual View

?

How to know this table is fragmented and where the fragments are?




Physical View


Challenge II: Global Catalog

- Centralized version (@master)
 - Accessing it is a bottleneck
 - Single-point failure
 - May add a mirror
 - Poorer performance
- Distributed version (several *masters*)
 - Replica synchronization
 - Potential inconsistencies

Challenge II: Global Catalog




Conceptual View



Physical View

Challenge II: Global Catalog



Conceptual View


Catalog:

T <<fragmentation strategy>>

F1: @S1, @S2, @S4


F2: @S2, @S3, @S4

F3: @S1, @S3, @S4, @Sn



Physical View

Challenge II: Global Catalog



Conceptual View

Catalog:


T <<fragmentation strategy>>

F1: @S1, @S2, @S4

F2: @S2, @S3, @S4

F3: @S1, @S3, @S4, @Sn

This information is
typically stored in a
distributed index




Physical View

Challenge II: Global Catalog

Conceptual View

Centralized Version




Physical View





Challenge II: Global Catalog

Conceptual View


Primary server



Centralized Version



...




Secondary servers

Physical View


Challenge II: Global Catalog

Conceptual View

Primary server



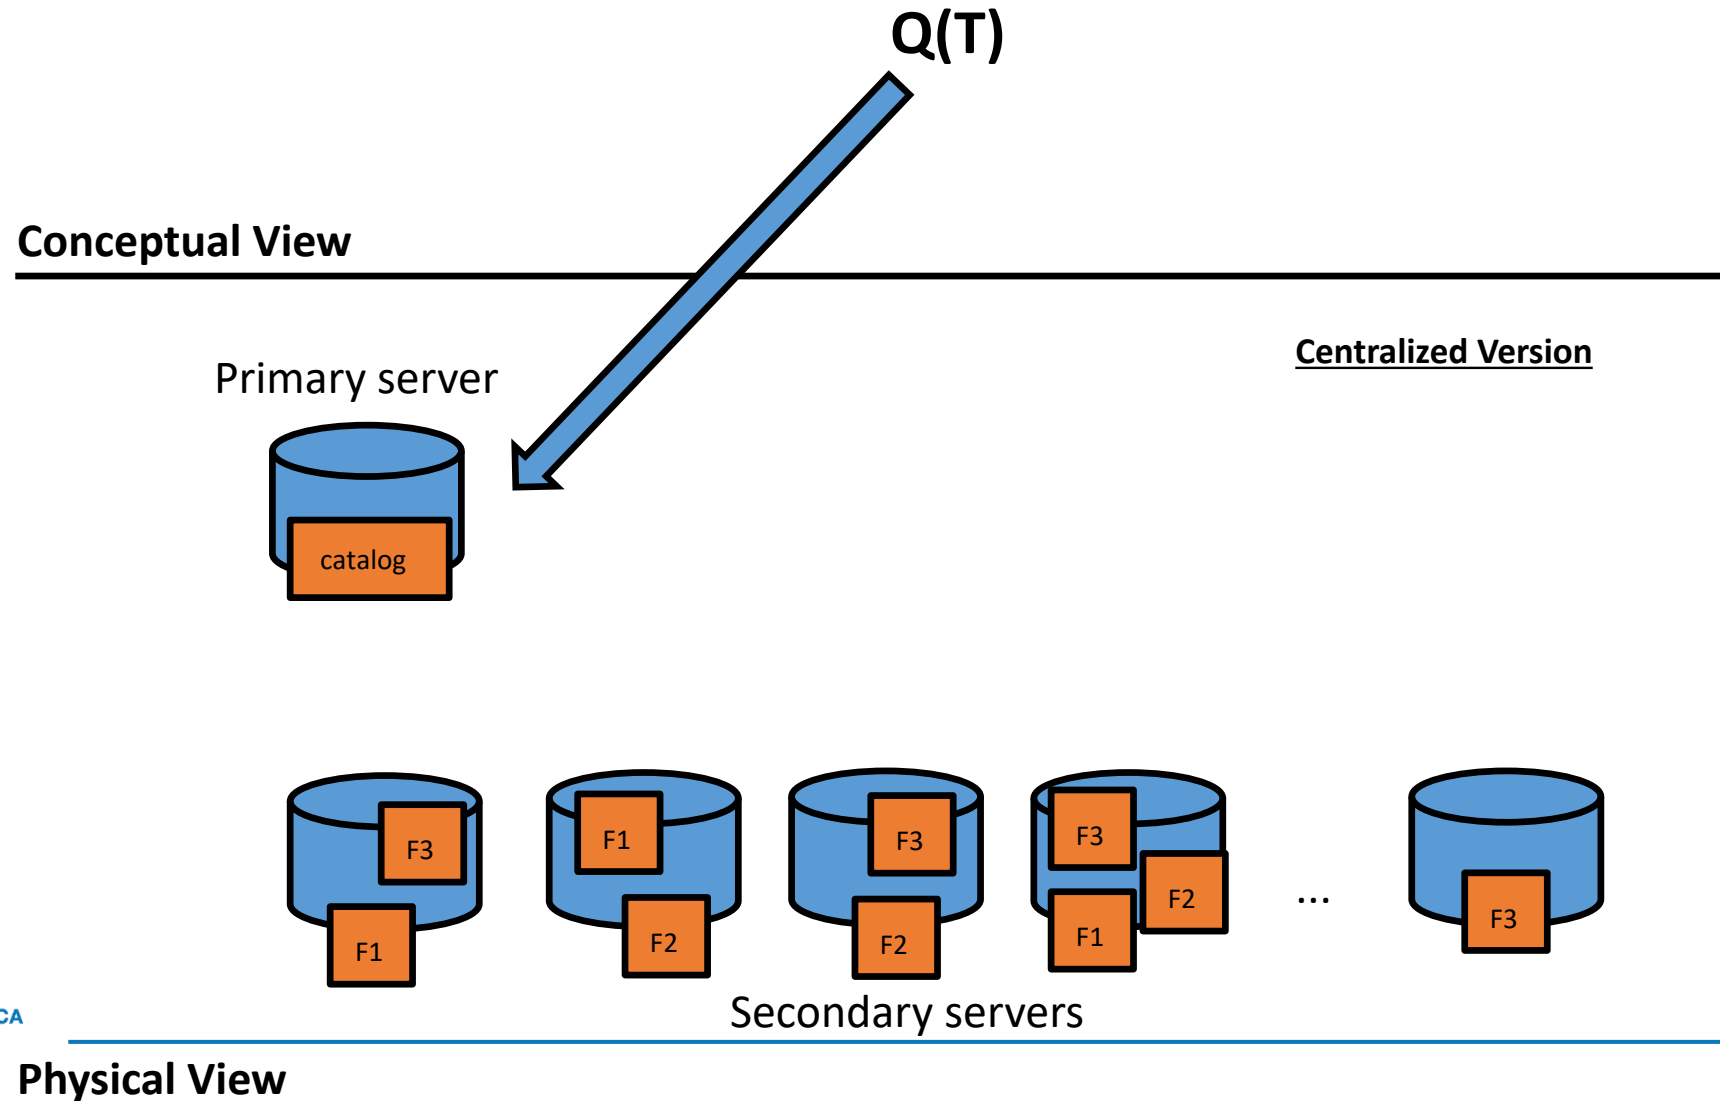
Centralized Version



Secondary servers

Physical View


Challenge II: Global Catalog







Challenge II: Global Catalog

Conceptual View


Primary server



Distributed Version



...




Secondary servers

Physical View





Challenge II: Global Catalog

Conceptual View


Primary server



Distributed Version



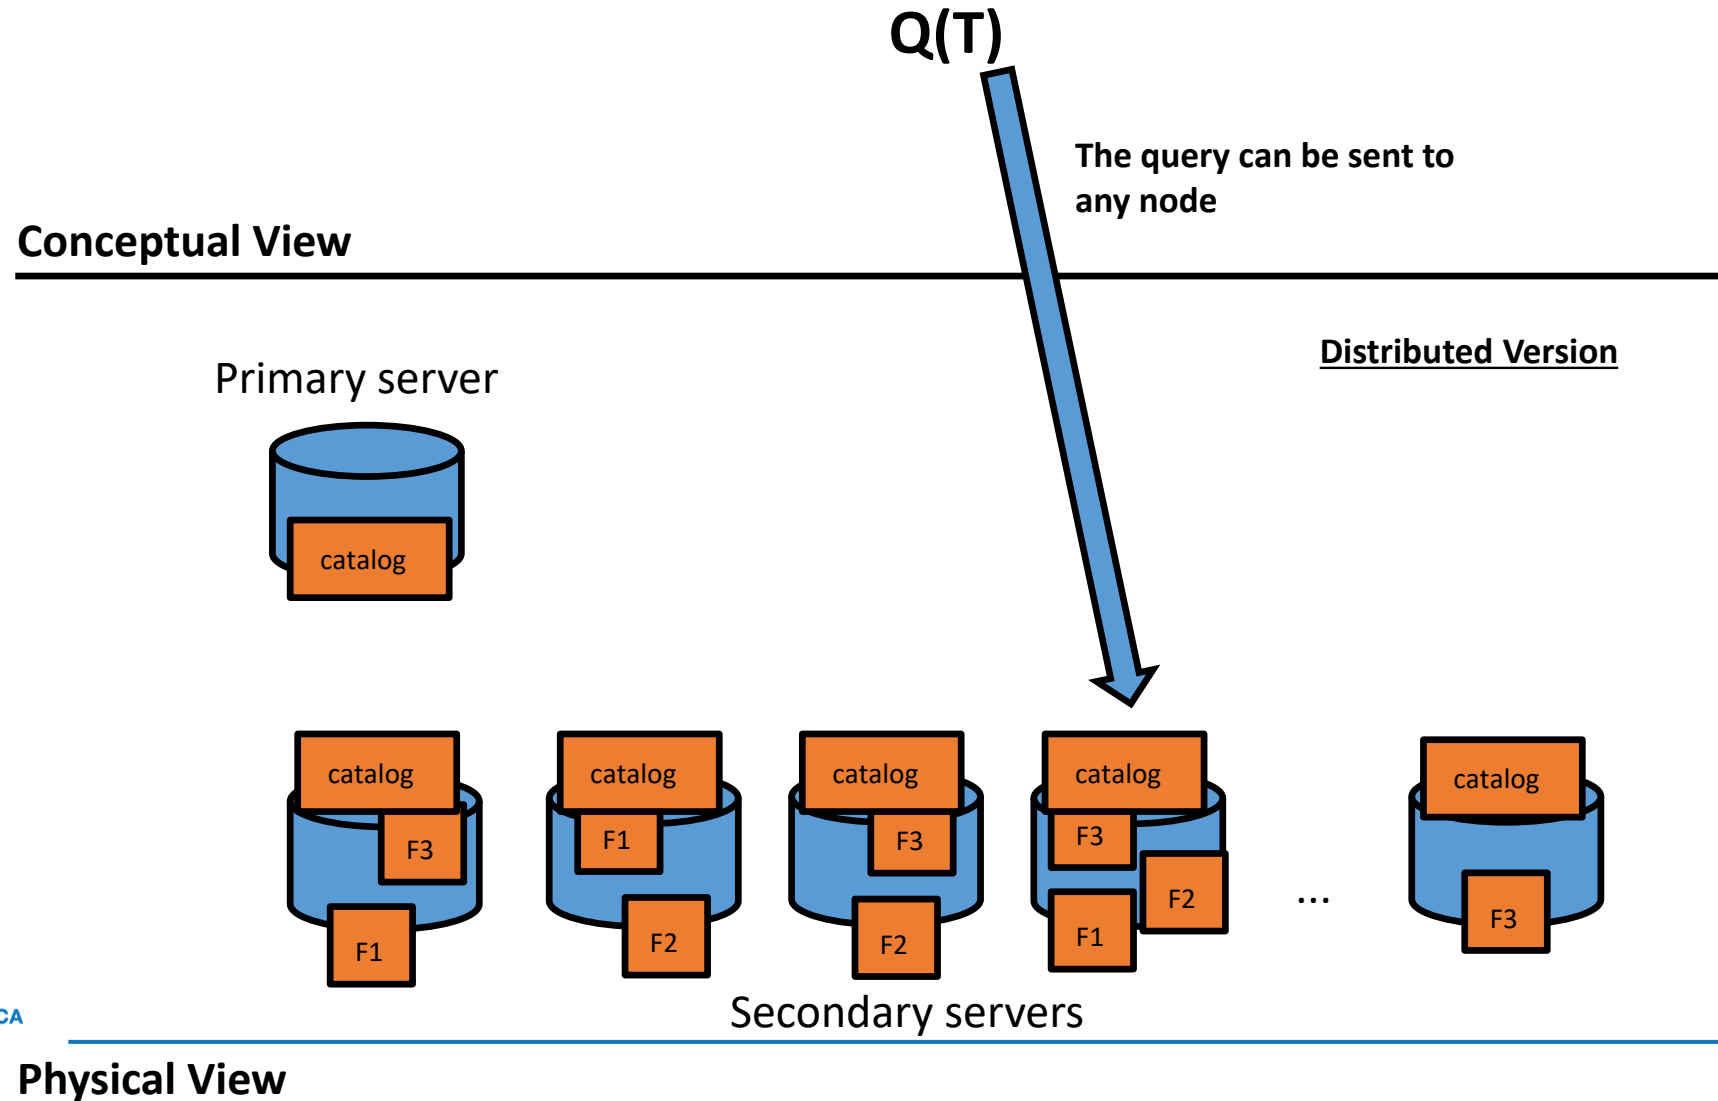
...



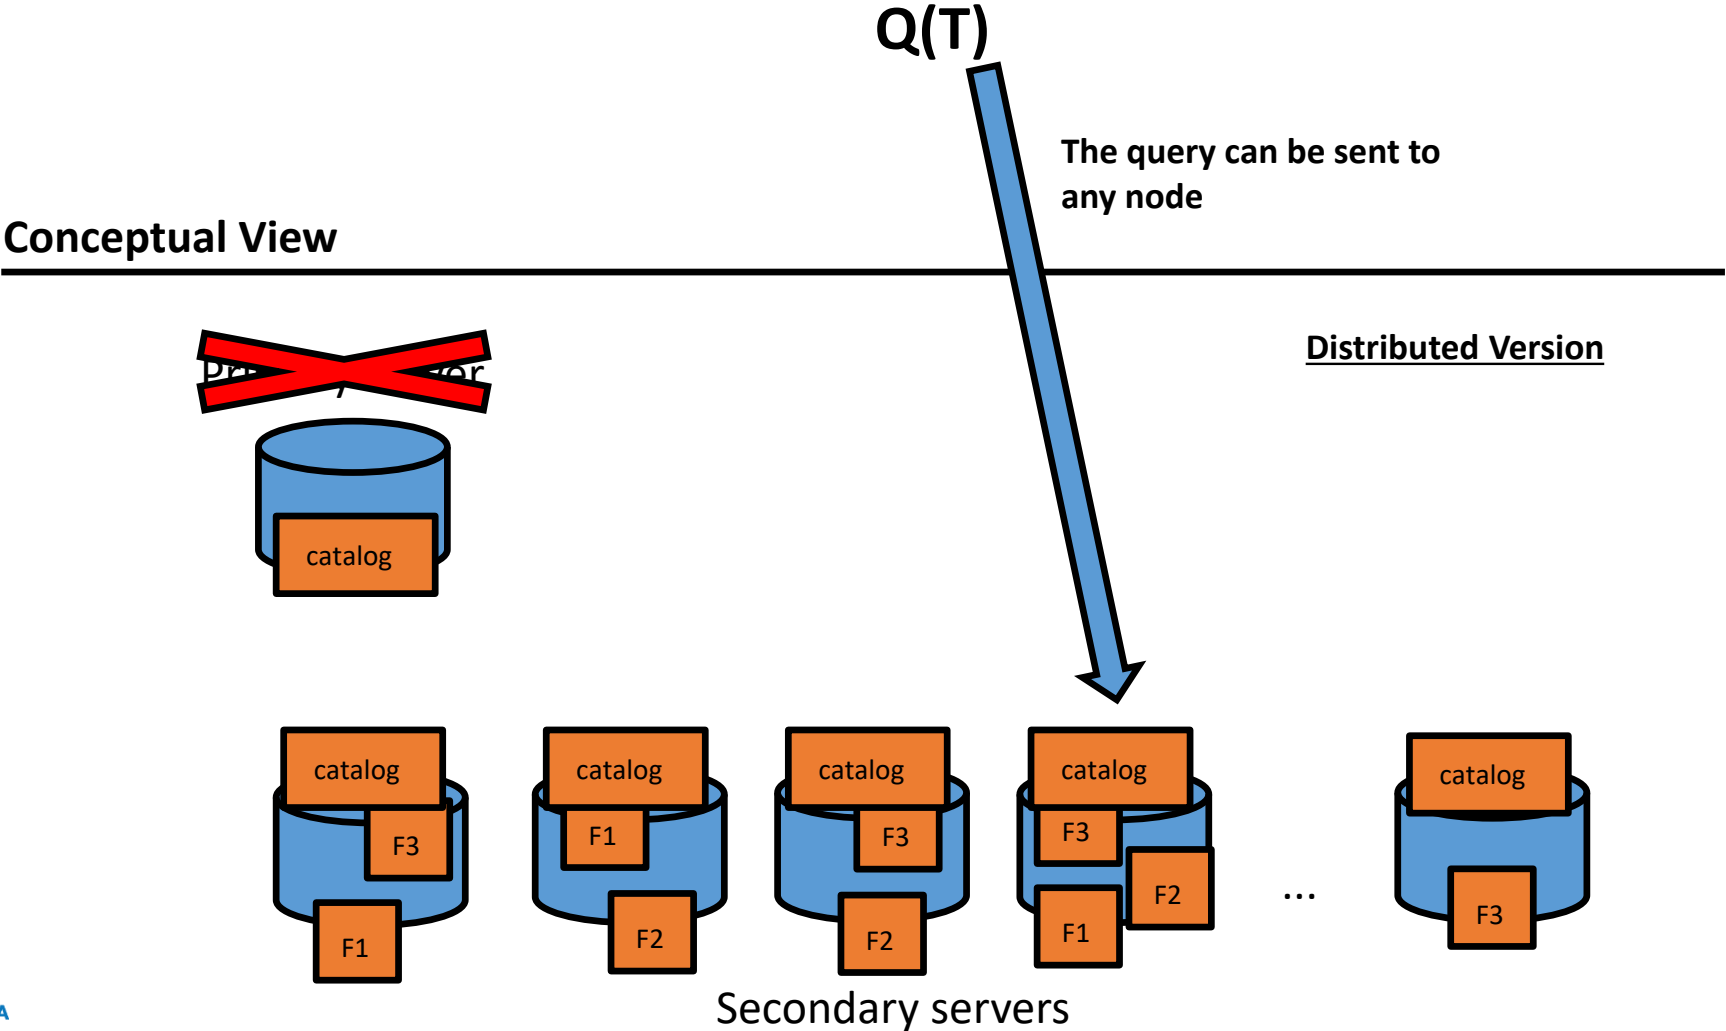
Secondary servers

Physical View

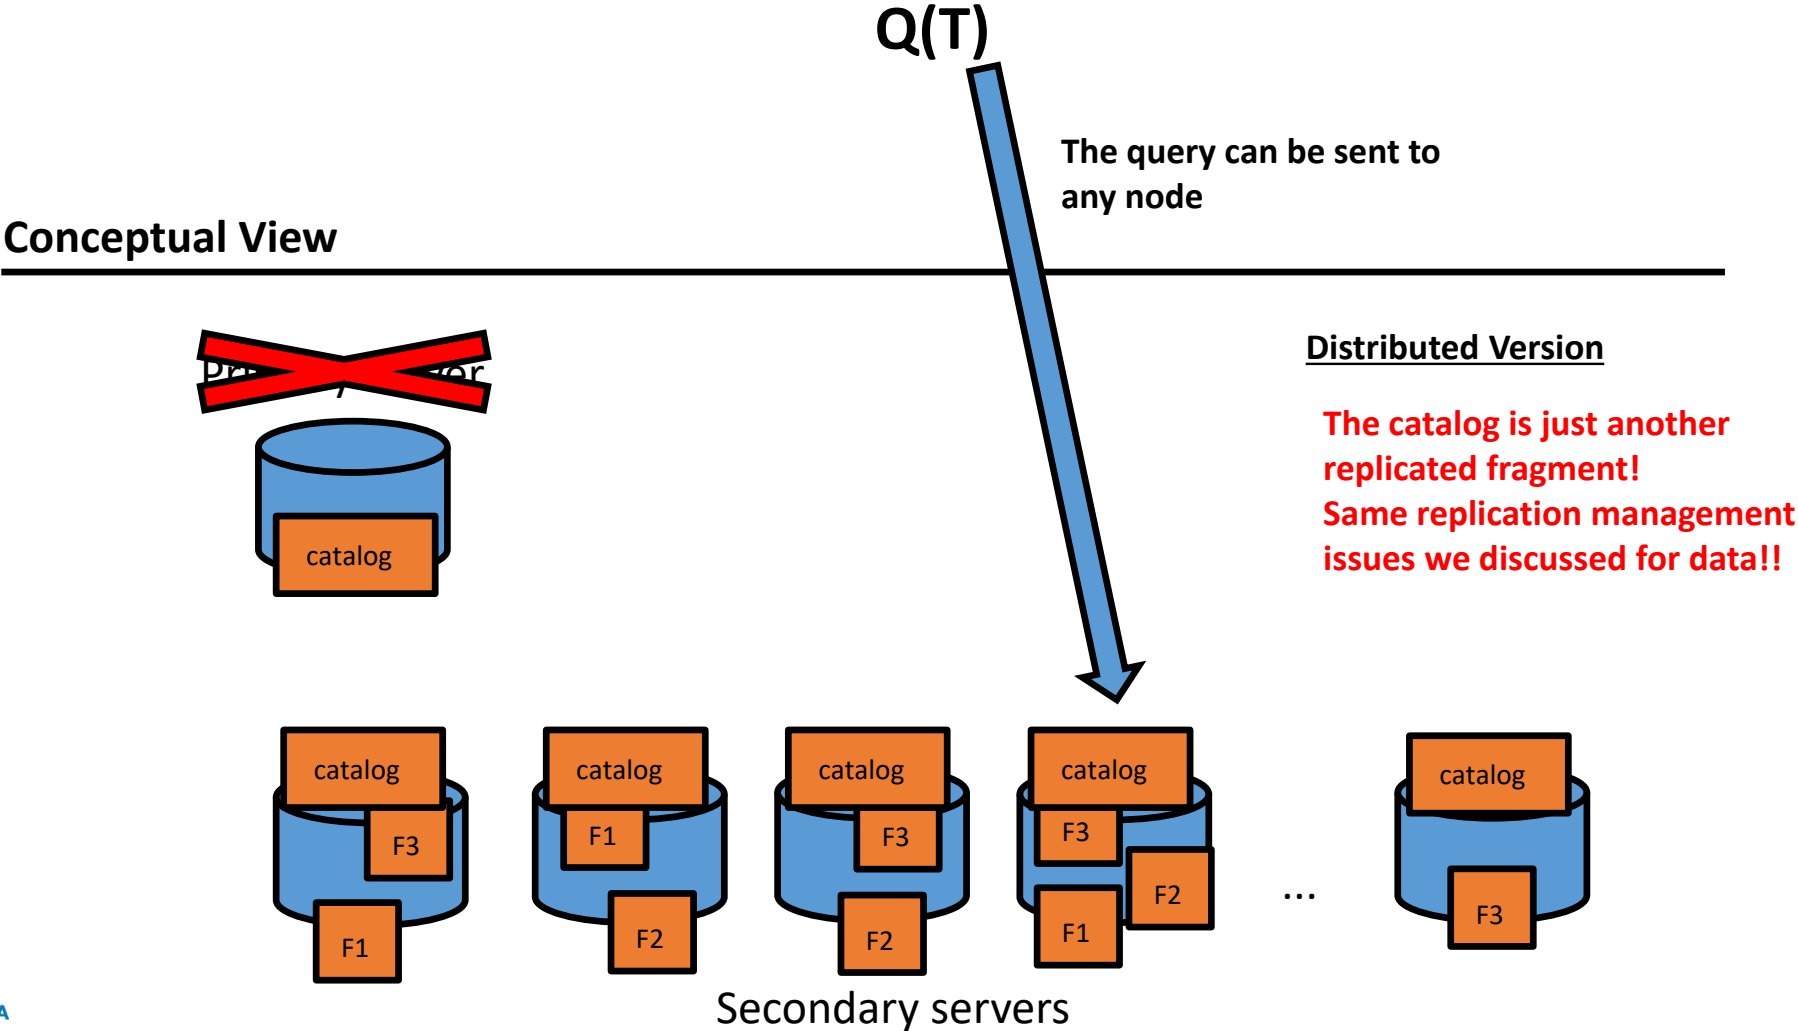
Challenge II: Global Catalog



Challenge II: Global Catalog



Challenge II: Global Catalog



Summary

- Benefits of distributed systems and their relevance for NOSQL
- What is a distributed DBMS
- System architecture of a DDBMS
 - Distribution transparency
 - Replication transparency
 - Fragmentation transparency
- Distributed Database design
 - Data fragmentation
 - Data allocation
 - Data replication
- Global catalog management

Bibliography

- M.T. Özsu and P. Valduriez. *Principles of Distributed Database Systems*. Second edition. Prentice Hall, 1999
- Serge Abiteboul, Ioana Manolescu, Philippe Rigaux, Marie-Christine Rousset, Pierre Senellart. *Web Data Management*. Cambridge Press, 2011.
- L. Liu, M.T. Özsu (Eds.). *Encyclopedia of Database Systems*. Springer, 2009